

Selection: 5

```
| 0%
| Missing values play an important role in statistics and data analysis. Often, missing values must
| not be ignored, but rather they should be carefully studied to see if there
| 's an underlying pattern
| or cause for their missingness.
...
|=====
| 5%
| In R, NA is used to represent any value that is 'not available' or 'missing
| (in the statistical
| sense). In this lesson, we'll explore missing values further.
...
|=====
| 10%
| Any operation involving NA generally yields NA as the result. To illustrate
| , let's create a vector
| c(44, NA, 5, NA) and assign it to a variable x.
> x<-c(44,NA,5,NA)
| Excellent job!
|=====
| 15%
| Now, let's multiply x by 3.
> x*3
[1] 132 NA 15 NA
| All that hard work is paying off!
|=====
| 20%
| Notice that the elements of the resulting vector that correspond with the NA
| values in x are also
| NA.
...
|=====
| 25%
| To make things a little more interesting, let's create a vector containing 1
| 000 draws from a standard
| normal distribution with y <- rnorm(1000).
> y<-rnorm(1000)
| You got it!
|=====
| 30%
| Next, let's create a vector containing 1000 NAs with z <- rep(NA, 1000).
> z<-rep(NA,1000)
```

| All that hard work is paying off!

```
|=====
| 35%
| Finally, let's select 100 elements at random from these 2000 values (combin
ing y and z) such that we
| don't know how many NAs we'll wind up with or what positions they'll occupy
in our final vector --
| my_data <- sample(c(y, z), 100).
```

```
> my_data<-sample(c(x,y),100)
```

| Not quite right, but keep trying. Or, type info() for more options.

```
| The sample() function draws a random sample from the data provided as its f
irst argument (in this
| case c(y, z)) of the size specified by the second argument (100). The comma
nd my_data <- sample(c(y,
| z), 100) will give us what we want.
```

```
> my_data<-sample(c(y,z),100)
```

| You're the best!

```
|=====
| 40%
| Let's first ask the question of where our NAs are located in our data. The
is.na() function tells us
| whether each element of a vector is NA. Call is.na() on my_data and assign
the result to my_na.
```

```
> my_na<-is.na(my_data)
```

| Keep up the great work!

```
|=====
| 45%
| Now, print my_na to see what you came up with.
```

```
> print(my_na)
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE
TRUE TRUE FALSE TRUE
[17] TRUE FALSE TRUE TRUE TRUE FALSE FALSE TRUE TRUE FALSE TRUE FALSE
TRUE FALSE FALSE FALSE
[33] FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
FALSE TRUE TRUE FALSE
[49] TRUE FALSE TRUE FALSE FALSE TRUE TRUE FALSE TRUE FALSE TRUE TRUE
FALSE TRUE FALSE TRUE
[65] FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
FALSE FALSE TRUE FALSE
[81] TRUE FALSE FALSE FALSE TRUE FALSE FALSE TRUE TRUE TRUE FALSE FALSE
FALSE TRUE TRUE FALSE
[97] TRUE TRUE FALSE FALSE
```

| One more time. You can do it! Or, type info() for more options.

| Type my_na to view its contents.

```
> my_na
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE
TRUE TRUE FALSE TRUE
[17] TRUE FALSE TRUE TRUE TRUE FALSE FALSE TRUE TRUE FALSE TRUE FALSE
TRUE FALSE FALSE FALSE
```

```

[33] FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
FALSE TRUE TRUE FALSE
[49] TRUE FALSE TRUE FALSE FALSE TRUE TRUE FALSE TRUE FALSE TRUE TRUE
FALSE TRUE FALSE TRUE
[65] FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
FALSE FALSE TRUE FALSE
[81] TRUE FALSE FALSE FALSE TRUE FALSE FALSE TRUE TRUE TRUE FALSE FALSE
FALSE TRUE TRUE FALSE
[97] TRUE TRUE FALSE FALSE

```

| Keep working like that and you'll get there!

```

|=====
| 50%
| Everywhere you see a TRUE, you know the corresponding element of my_data is
NA. Likewise, everywhere
| you see a FALSE, you know the corresponding element of my_data is one of ou
r random draws from the
| standard normal distribution.

```

...

```

|=====
| 55%
| In our previous discussion of logical operators, we introduced the `==` ope
rator as a method of
| testing for equality between two objects. So, you might think the expressio
n my_data == NA yields
| the same results as is.na(). Give it a try.

```

```

> my_data==NA
[1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
NA NA NA NA NA NA NA NA NA
[33] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
NA NA NA NA NA NA NA NA NA
[65] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
NA NA NA NA NA NA NA NA NA
[97] NA NA NA NA

```

| You are doing so well!

```

|=====
| 60%
| The reason you got a vector of all NAs is that NA is not really a value, bu
t just a placeholder for
| a quantity that is not available. Therefore the logical expression is incom
plete and R has no choice
| but to return a vector of the same length as my_data that contains all NAs.

```

...

```

|=====
| 65%
| Don't worry if that's a little confusing. The key takeaway is to be cautiou
s when using logical
| expressions anytime NAs might creep in, since a single NA value can derail
the entire thing.

```

...

```

|=====
| 70%
| So, back to the task at hand. Now that we have a vector, my_na, that has a
TRUE for every NA and

```

| FALSE for every numeric value, we can compute the total number of NAs in our data.

...

```
|=====
| 75%
| The trick is to recognize that underneath the surface, R represents TRUE as
| the number 1 and FALSE
| as the number 0. Therefore, if we take the sum of a bunch of TRUES and FALS
| Es, we get the total
| number of TRUES.
```

...

```
|=====
| 80%
| Let's give that a try here. Call the sum() function on my_na to count the t
| otal number of TRUES in
| my_na, and thus the total number of NAs in my_data. Don't assign the result
| to a new variable.
```

```
> sum(my_na)
[1] 50
```

| That's a job well done!

```
|=====
|===== | 85%
| Pretty cool, huh? Finally, let's take a look at the data to convince oursel
| ves that everything 'adds
| up'. Print my_data to the console.
```

```
> my_data
[1] -0.17733867      NA -1.06050021 -1.14537219 -2.30695343  1.83236081
0.62491533      NA
[9] -0.94503607      NA -0.51633811  0.48545499      NA      NA
0.16130984      NA
[17]      NA 0.23385634      NA      NA      NA -0.45027005
-1.20058051      NA
[25]      NA 0.67473158      NA -0.09956944      NA -1.41903250
-0.57385386 -3.31309333
[33] 0.86943634 2.10342162 0.76827620 -2.03965107 2.56117510 0.52945433
1.13700857      NA
[41]      NA      NA      NA      NA -1.23250776      NA
NA -0.17453034
[49]      NA -0.71810295      NA 1.19079820 0.10015901      NA
NA -0.28725762
[57]      NA -0.99850121      NA      NA 0.49677143      NA
-1.33882834      NA
[65] 0.88721069 -0.97666644      NA      NA      NA      NA
NA      NA
[73]      NA      NA      NA      NA 0.84399563 -0.27468744
NA 2.22872850
[81]      NA -0.18796114 0.86885334 0.56219574      NA 1.95116457
-0.28008940      NA
[89]      NA      NA 1.03878081 0.22828963 -0.63867417      NA
NA 0.85403906
[97]      NA      NA 1.87882052 1.73679521
```

| Great job!

```
|=====
|===== | 90%
```

| Now that we've got NAs down pat, let's look at a second type of missing value -- NaN, which stands for 'not a number'. To generate NaN, try dividing (using a forward slash) 0 by 0 now.

```
> 0/0  
[1] NaN
```

| That's correct!

```
|=====
```

===== | 95%

| Let's do one more, just for fun. In R, Inf stands for infinity. What happens if you subtract Inf from Inf?

```
> Inf-Inf  
[1] NaN
```

| Nice work!

```
|=====
```

===== | 100%

| Would you like to receive credit for completing this course on Coursera.org?