

Selection: 1

| Please choose a lesson, or type 0 to return to course menu.

1: Basic Building Blocks Numbers	2: Workspace and Files	3: Sequences of Numbers
4: Vectors ctors	5: Missing Values	6: Subsetting ve ctors
7: Matrices and Data Frames	8: Logic	9: Functions
10: lapply and sapply	11: vapply and tapply	12: Looking at Da ta
13: Simulation	14: Dates and Times	15: Base Graphics

Selection: 6

|
| 0%

| In this lesson, we'll see how to extract elements from a vector based on some conditions that we specify.

...

|==
| 3%

| For example, we may only be interested in the first 20 elements of a vector, or only the elements that are not NA, or only those that are positive or correspond to a specific variable of interest.
| By the end of this lesson, you'll know how to handle each of these scenarios.

...

|=====
| 5%

| I've created for you a vector called x that contains a random ordering of 20 numbers (from a standard normal distribution) and 20 NAs. Type x now to see what it looks like.

```
> x
[1] NA NA NA NA NA -0.94887718
NA 0.13725073
[9] NA NA -1.85099548 NA 0.94707769 0.47929006
-0.23134894 NA
[17] 0.03379646 1.41590243 NA -0.37050224 0.05623335 1.19070883
-1.19223571 1.56108480
[25] NA -0.85689449 0.32979317 0.26235968 -1.38318148 NA
-0.43592424 -0.28931595
[33] NA -0.26026676 NA NA NA NA
NA NA
```

| That's correct!

|=====
| 8%

| The way you tell R that you want to select some particular elements (i.e. a 'subset') from a vector is by placing an 'index vector' in square brackets immediately following the name of the vector.

...

|=====

| 10%

| For a simple example, try x[1:10] to view the first ten elements of x.

> x[1:10]

```
[1] NA NA NA NA NA -0.9488772
NA 0.1372507 NA
[10] NA
```

| You are amazing!

|=====

| 13%

| Index vectors come in four different flavors -- logical vectors, vectors of positive integers, vectors of negative integers, and vectors of character strings -- each of which we'll cover in this lesson.

...

|=====

| 15%

| Let's start by indexing with logical vectors. One common scenario when working with real-world data is that we want to extract all elements of a vector that are not NA (i.e. missing data). Recall that is.na(x) yields a vector of logical values the same length as x, with TRUEs corresponding to NA values in x and FALSEs corresponding to non-NA values in x.

...

|=====

| 18%

| What do you think x[is.na(x)] will give you?

- 1: A vector of length 0
- 2: A vector of all NAs
- 3: A vector of TRUEs and FALSEs
- 4: A vector with no NAs

selection: 3

| That's not exactly what I'm looking for. Try again.

| Remember that is.na(x) tells us where the NAs are in a vector. So if we subset x based on that, what do you expect to happen?

- 1: A vector of all NAs
- 2: A vector of length 0
- 3: A vector with no NAs
- 4: A vector of TRUEs and FALSEs

selection: 2

| One more time. You can do it!

| Remember that is.na(x) tells us where the NAs are in a vector. So if we subset x based on that, what do you expect to happen?

- 1: A vector of TRUEs and FALSEs
- 2: A vector with no NAs
- 3: A vector of all NAs
- 4: A vector of length 0

Selection: 1

| Keep trying!

| Remember that `is.na(x)` tells us where the NAs are in a vector. So if we subset `x` based on that, what do you expect to happen?

- 1: A vector with no NAs
- 2: A vector of all NAs
- 3: A vector of TRUEs and FALSEs
- 4: A vector of length 0

Selection: 2

| Excellent job!

| =====
| 21%

| Prove it to yourself by typing `x[is.na(x)]`.

> `x[is.na(x)]`

```
[1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
```

| You are doing so well!

| =====
| 23%

| Recall that `!` gives us the negation of a logical expression, so `!is.na(x)` can be read as 'is not NA'. Therefore, if we want to create a vector called `y` that contains all of the non-NA values from `x`, we can use `y <- x[!is.na(x)]`. Give it a try.

> `y <- x[!is.na(x)]`

| Perseverance, that's the answer.

| =====
| 26%

| Print `y` to the console.

> `y`

```
[1] -0.94887718  0.13725073 -1.85099548  0.94707769  0.47929006 -0.23134894  
0.03379646  1.41590243  
[9] -0.37050224  0.05623335  1.19070883 -1.19223571  1.56108480 -0.85689449  
0.32979317  0.26235968  
[17] -1.38318148 -0.43592424 -0.28931595 -0.26026676
```

| That's correct!

| =====
| 28%

| Now that we've isolated the non-missing values of `x` and put them in `y`, we can subset `y` as we please.

...

```
|=====
| 31%
| Recall that the expression y > 0 will give us a vector of logical values the
| same length as y, with
| TRUES corresponding to values of y that are greater than zero and FALSEs co
| rresponding to values of
| y that are less than or equal to zero. What do you think y[y > 0] will give
| you?
```

- 1: A vector of all the negative elements of y
- 2: A vector of length 0
- 3: A vector of TRUES and FALSEs
- 4: A vector of all NAs
- 5: A vector of all the positive elements of y

Selection: 5

| Nice work!

```
|=====
| 33%
| Type y[y > 0] to see that we get all of the positive elements of y, which a
| re also the positive
| elements of our original vector x.
```

```
> y[y>0]
[1] 0.13725073 0.94707769 0.47929006 0.03379646 1.41590243 0.05623335 1.1907
0883 1.56108480 0.32979317
[10] 0.26235968
```

| That's correct!

```
|=====
| 36%
| You might wonder why we didn't just start with x[x > 0] to isolate the posi
| tive elements of x. Try
| that now to see why.
```

```
> x[x>0]
[1] NA NA NA NA NA NA NA 0.1372
5073 NA NA
[10] NA 0.94707769 0.47929006 NA 0.03379646 1.41590243
NA 0.05623335 1.19070883
[19] 1.56108480 NA 0.32979317 0.26235968 NA NA
NA NA NA
[28] NA NA NA
```

| You are doing so well!

```
|=====
| 38%
| Since NA is not a value, but rather a placeholder for an unknown quantity,
| the expression NA > 0
| evaluates to NA. Hence we get a bunch of NAs mixed in with our positive num
| bers when we do this.
```

...

```
|=====
| 41%
| Combining our knowledge of logical operators with our new knowledge of subs
| etting, we could do this
| -- x[!is.na(x) & x > 0]. Try it out.
```

```
> x[!is.na(x)$x>0]
Error in is.na(x)$x : $ operator is invalid for atomic vectors
> x[!is.na(x)&x>0]
[1] 0.13725073 0.94707769 0.47929006 0.03379646 1.41590243 0.05623335 1.1907
0883 1.56108480 0.32979317
[10] 0.26235968
```

| Keep working like that and you'll get there!

```
|=====
| 44%
| In this case, we request only values of x that are both non-missing AND gre
ater than zero.
```

...

```
|=====
| 46%
| I've already shown you how to subset just the first ten values of x using x
[1:10]. In this case,
| we're providing a vector of positive integers inside of the square brackets
, which tells R to return
| only the elements of x numbered 1 through 10.
```

...

```
|=====
| 49%
| Many programming languages use what's called 'zero-based indexing', which m
eans that the first
| element of a vector is considered element 0. R uses 'one-based indexing', w
hich (you guessed it!)
| means the first element of a vector is considered element 1.
```

...

```
|=====
| 51%
| Can you figure out how we'd subset the 3rd, 5th, and 7th elements of x? Hin
t -- Use the c() function
| to specify the element numbers as a numeric vector.
```

```
> x[c(3,5,7)]
[1] NA NA NA
```

| Nice work!

```
|=====
| 54%
| It's important that when using integer vectors to subset our vector x, we s
tick with the set of
| indexes {1, 2, ..., 40} since x only has 40 elements. What happens if we as
k for the zeroth element
| of x (i.e. x[0])? Give it a try.
```

```
> x[0]
numeric(0)
```

| You are doing so well!

```
|=====
| 56%
| As you might expect, we get nothing useful. Unfortunately, R doesn't preven
t us from doing this.
```

| what if we ask for the 3000th element of x? Try it out.

```
> x[3000]
[1] NA
```

| keep working like that and you'll get there!

```
|=====
| 59%
| Again, nothing useful, but R doesn't prevent us from asking for it. This sh
ould be a cautionary
| tale. You should always make sure that what you are asking for is within th
e bounds of the vector
| you're working with.
```

...

```
|=====
| 62%
| what if we're interested in all elements of x EXCEPT the 2nd and 10th? It w
ould be pretty tedious to
| construct a vector containing all numbers 1 through 40 EXCEPT 2 and 10.
```

...

```
|=====
| 64%
| Luckily, R accepts negative integer indexes. Whereas x[c(2, 10)] gives us O
NLY the 2nd and 10th
| elements of x, x[c(-2, -10)] gives us all elements of x EXCEPT for the 2nd
and 10 elements. Try
| x[c(-2, -10)] now to see this.
```

```
> x[c(-2, -10)]
[1] NA NA NA NA -0.94887718 NA
0.13725073 NA
[9] -1.85099548 NA 0.94707769 0.47929006 -0.23134894 NA
0.03379646 1.41590243
[17] NA -0.37050224 0.05623335 1.19070883 -1.19223571 1.56108480
NA -0.85689449
[25] 0.32979317 0.26235968 -1.38318148 NA -0.43592424 -0.28931595
NA -0.26026676
[33] NA NA NA NA NA NA NA
```

| Excellent job!

```
|=====
| 67%
| A shorthand way of specifying multiple negative numbers is to put the negat
ive sign out in front of
| the vector of positive numbers. Type x[-c(2, 10)] to get the exact same res
ult.
```

```
> x[-c(2, 10)]
[1] NA NA NA NA -0.94887718 NA
0.13725073 NA
[9] -1.85099548 NA 0.94707769 0.47929006 -0.23134894 NA
0.03379646 1.41590243
[17] NA -0.37050224 0.05623335 1.19070883 -1.19223571 1.56108480
NA -0.85689449
[25] 0.32979317 0.26235968 -1.38318148 NA -0.43592424 -0.28931595
NA -0.26026676
[33] NA NA NA NA NA NA NA
```

| Great job!

```
|=====
| 69%
| So far, we've covered three types of index vectors -- logical, positive integer, and negative integer. The only remaining type requires us to introduce the concept of 'named' elements.
```

...

```
|=====
| 72%
| Create a numeric vector with three named elements using vect <- c(foo = 11, bar = 2, norf = NA).
```

```
> vect<-c(foo=11,bar=2,norf=NA)
```

| Great job!

```
|=====
| 74%
| When we print vect to the console, you'll see that each element has a name. Try it out.
```

```
> vect
foo bar norf
11  2  NA
```

| Excellent job!

```
|=====
| 77%
| We can also get the names of vect by passing vect as an argument to the names() function. Give that a try.
```

```
> names(vect)
[1] "foo" "bar" "norf"
```

| You nailed it! Good job!

```
|=====
| 79%
| Alternatively, we can create an unnamed vector vect2 with c(11, 2, NA). Do that now.
```

```
> vect<-c(11,2,NA)
```

| Try again. Getting it right on the first try is boring anyway! Or, type info() for more options.

| Create an ordinary (unnamed) vector called vect2 that contains c(11, 2, NA).

```
> vect2<-c(11,2,NA)
```

| Your dedication is inspiring!

```
|=====
== | 82%
| Then, we can add the `names` attribute to vect2 after the fact with names(vect2) <- c("foo", "bar", "norf"). Go ahead.
```

```
> names(vect2) <- c("foo", "bar","norf")
```

```
| Nice work!
```

```
|=====
===== | 85%
| Now, let's check that vect and vect2 are the same by passing them as arguments to the identical() function.
```

```
> identical(vect,vect2)
[1] TRUE
```

```
| Great job!
```

```
|=====
===== | 87%
| Indeed, vect and vect2 are identical named vectors.
```

```
...
```

```
|=====
===== | 90%
| Now, back to the matter of subsetting a vector by named elements. Which of the following commands do you think would give us the second element of vect?
```

```
1: vect["bar"]
2: vect["2"]
3: vect[bar]
```

```
selection: 1
```

```
| You are really on a roll!
```

```
|=====
===== | 92%
| Now, try it out.
```

```
> vec[t"bar"]
Error: unexpected string constant in "vec[t"bar""
> vec["bar"]
Error: object 'vec' not found
> vect["bar"]
bar
2
```

```
| Excellent work!
```

```
|=====
===== | 95%
| Likewise, we can specify a vector of names with vect[c("foo", "bar")]. Try it out.
```

```
> vect[c("foo", "bar")]
foo bar
11 2
```

```
| Perseverance, that's the answer.
```

```
|=====
===== | 97%
```


| Now you know all four methods of subsetting data from vectors. Different approaches are best in different scenarios and when in doubt, try it out!

...

|=====|
=====| 100%
| would you like to receive credit for completing this course on Coursera.org?
?

- 1: No
- 2: Yes