

Selection: 14

```
|
| 0%

| R has a special way of representing dates and times, which can be helpful i
f you're working with
| data that show how something changes over time (i.e. time-series data) or i
f your data contain some
| other temporal information, like dates of birth.

...

|===
| 3%
| Dates are represented by the 'Date' class and times are represented by the
'POSIXct' and 'POSIXlt'
| classes. Internally, dates are stored as the number of days since 1970-01-0
1 and times are stored as
| either the number of seconds since 1970-01-01 (for 'POSIXct') or a list of
seconds, minutes, hours,
| etc. (for 'POSIXlt').

...

|=====
| 6%
| Let's start by using d1 <- Sys.Date() to get the current date and store it
in the variable d1.
| (That's the letter 'd' and the number 1.)

> d1 <- Sys.Date()

| Keep up the great work!

|=====
| 8%
| Use the class() function to confirm d1 is a Date object.

> class(d1)
[1] "Date"

| Keep up the great work!

|=====
| 11%
| We can use the unclass() function to see what d1 looks like internally. Try
it out.

> unclass(d1)
[1] 18368

| You are doing so well!

|=====
| 14%
| That's the exact number of days since 1970-01-01!

...

|=====
| 17%
| However, if you print d1 to the console, you'll get today's date -- YEAR-MO
NTH-DAY. Give it a try.
```

```

> d1
[1] "2020-04-16"

| You are doing so well!

|=====
| 19%
| What if we need to reference a date prior to 1970-01-01? Create a variable
d2 containing
| as.Date("1969-01-01").

> d2<-as.Date("1969-01-01")

| You're the best!

|=====
| 22%
| Now use unclass() again to see what d2 looks like internally.

> unclass(d2)
[1] -365

| You got it!

|=====
| 25%
| As you may have anticipated, you get a negative number. In this case, it's
-365, since 1969-01-01 is
| exactly one calendar year (i.e. 365 days) BEFORE 1970-01-01.

...

|=====
| 28%
| Now, let's take a look at how R stores times. You can access the current da
te and time using the
| Sys.time() function with no arguments. Do this and store the result in a va
riable called t1.

> t1<- Sys.time()

| Almost! Try again. Or, type info() for more options.

| t1 <- Sys.time() will store the current date and time in a variable called
t1.

> t1<- Sys.time()

| You're the best!

|=====
| 31%
| View the contents of t1.

> t1
[1] "2020-04-16 22:39:16 IST"

| You got it right!

|=====
| 33%
| And check the class() of t1.

```

```
> class(t1)
[1] "POSIXct" "POSIXt"

| Excellent job!

|=====
| 36%
| As mentioned earlier, POSIXct is just one of two ways that R represents time information. (You can ignore the second value above, POSIXt, which just functions as a common language between POSIXct and POSIXlt.) Use unclass() to see what t1 looks like internally -- the (large) number of seconds since the beginning of 1970.
```

```
> unclass(t1)
[1] 1587056957

| Your dedication is inspiring!

|=====
| 39%
| By default, Sys.time() returns an object of class POSIXct, but we can coerce the result to POSIXlt with as.POSIXlt(Sys.time()). Give it a try and store the result in t2.
```

```
> t2<-as.POSIXlt(Sys.time())

| You nailed it! Good job!

|=====
| 42%
| Check the class of t2.
```

```
> t2
[1] "2020-04-16 22:41:35 IST"

| Nice try, but that's not exactly what I was hoping for. Try again. Or, type info() for more options.

| Type class(t2) to view its class.
```

```
> class(t2)
[1] "POSIXlt" "POSIXt"

| That's the answer I was looking for.

|=====
| 44%
| Now view its contents.
```

```
> t2
[1] "2020-04-16 22:41:35 IST"

| Keep working like that and you'll get there!

|=====
| 47%
| The printed format of t2 is identical to that of t1. Now unclass() t2 to see how it is different internally.
```

```
> unclass(t2)
$sec
```

```
[1] 35.57194
```

```
$min  
[1] 41
```

```
$hour  
[1] 22
```

```
$mday  
[1] 16
```

```
$mon  
[1] 3
```

```
$year  
[1] 120
```

```
$wday  
[1] 4
```

```
$yday  
[1] 106
```

```
$isdst  
[1] 0
```

```
$zone  
[1] "IST"
```

```
$gmtoff  
[1] 19800
```

```
attr(,"tzone")  
[1] "" "IST" "+0630"
```

| All that hard work is paying off!

```
|=====
```

| 50%

| t2, like all POSIXlt objects, is just a list of values that make up the date and time. Use

| str(unclass(t2)) to have a more compact view.

```
> str(unclass(t2))  
List of 11  
 $ sec   : num 35.6  
 $ min   : int 41  
 $ hour  : int 22  
 $ mday  : int 16  
 $ mon   : int 3  
 $ year  : int 120  
 $ wday  : int 4  
 $ yday  : int 106  
 $ isdst : int 0  
 $ zone  : chr "IST"  
 $ gmtoff: int 19800  
- attr(*, "tzone")= chr [1:3] "" "IST" "+0630"
```

| Keep working like that and you'll get there!

```
|=====
```

| 53%

| If, for example, we want just the minutes from the time stored in t2, we can access them with

| t2\$min. Give it a try.

```
> t2$min  
[1] 41
```

| You are amazing!

|=====

| 56%

| Now that we have explored all three types of date and time objects, let's look at a few functions that extract useful information from any of these objects -- weekdays(), months(), and quarters().

...

|=====

| 58%

| The weekdays() function will return the day of week from any date or time object. Try it out on d1, which is the Date object that contains today's date.

```
> weekdays(d1)  
[1] "Thursday"
```

| You got it right!

|=====

| 61%

| The months() function also works on any date or time object. Try it on t1, which is the POSIXct object that contains the current time (well, it was the current time when you created it).

```
> months(t1)  
[1] "April"
```

| You nailed it! Good job!

|=====

| 64%

| The quarters() function returns the quarter of the year (Q1-Q4) from any date or time object. Try it on t2, which is the POSIXlt object that contains the time at which you created it.

```
> quarters(t2)  
[1] "Q2"
```

| All that hard work is paying off!

|=====

| 67%

| Often, the dates and times in a dataset will be in a format that R does not recognize. The strptime() function can be helpful in this situation.

...

|=====

| 69%

| strptime() converts character vectors to POSIXlt. In that sense, it is similar to as.POSIXlt(),

| except that the input doesn't have to be in a particular format (YYYY-MM-DD
|).

...

|=====|
| 72%
| To see how it works, store the following character string in a variable called t3: "October 17, 1986 08:24" (with the quotes).

> t3<-"October 17, 1986 08:24"

| That's a job well done!

|=====|
| 75%
| Now, use strptime(t3, "%B %d, %Y %H:%M") to help R convert our date/time object to a format that it understands. Assign the result to a new variable called t4. (You should pull up the documentation for strptime() if you'd like to know more about how it works.)

> t4<-strptime(t3, "%B %d, %Y %H:%M")

| You got it right!

|=====|
| 78%
| Print the contents of t4.

> t4
[1] "1986-10-17 08:24:00 IST"

| Keep up the great work!

|=====|
| 81%
| That's the format we've come to expect. Now, let's check its class().

> class(t4)
[1] "POSIXlt" "POSIXt"

| Keep up the great work!

|=====|
| 83%
| Finally, there are a number of operations that you can perform on dates and times, including arithmetic operations (+ and -) and comparisons (<, ==, etc.)

...

|=====|
| 86%
| The variable t1 contains the time at which you created it (recall you used Sys.time()). Confirm that some time has passed since you created t1 by using the 'greater than' operator to compare it to the current time: Sys.time() > t1

> Sys.time() > t1
[1] TRUE

| Nice work!

```
|=====
=====| 89%
| So we know that some time has passed, but how much? Try subtracting t1 from
the current time using
| Sys.time() - t1. Don't forget the parentheses at the end of Sys.time(), sin
ce it is a function.
```

```
> Sys.time() - t1
Time difference of 13.66982 mins
```

| You are doing so well!

```
|=====
=====| 92%
| The same line of thinking applies to addition and the other comparison oper
ators. If you want more
| control over the units when finding the above difference in times, you can
use difftime(), which
| allows you to specify a 'units' parameter.
```

...

```
|=====
=====| 94%
| Use difftime(Sys.time(), t1, units = 'days') to find the amount of time in
DAYS that has passed
| since you created t1.
```

```
> difftime(Sys.time(), t1, units = 'days')
Time difference of 0.0100779 days
```

| You are doing so well!

```
|=====
=====| 97%
| In this lesson, you learned how to work with dates and times in R. While it
is important to
| understand the basics, if you find yourself working with dates and times of
ten, you may want to
| check out the lubridate package by Hadley Wickham.
```

...

```
|=====
=====| 100%
| Would you like to receive credit for completing this course on Coursera.org
?
```